

# IN.5022 — Concurrent and Distributed Computing Foundations

Prof. P. Felber

[pascal.felber@unine.ch](mailto:pascal.felber@unine.ch)

# Agenda



- What is a distributed system?
- Architectural models
- Fundamental models

# What is a distributed system?



A collection of independent computers that appears to its users as a ***single coherent system***

Provides the means for performance, scalability, dependability, etc.

# What is a distributed system?

A collection of applications communicating by message passing in order to solve ***a common task***

Independent computers cooperate together

# What is a distributed system?

“A distributed system is one in which the *failure* of a computer you didn’t even know existed can render your own computer unusable”

Leslie Lamport




Introduces special problems regarding correctness, complexity, failures...

**Fault tolerance:** the ability of a system to provide useful service (possibly degraded in functionality and/or performance), despite the fact that some of its components malfunction

# What is a distributed system?

- Abstract view: a **network** of **processes**
  - *Nodes* are processes, *edges* are communication channels




# Example of distributed systems

- Internet, intranets, WWW
- Telecommunication networks
- Airline reservation systems
- Aircraft control systems
- Electronic banking (interbank networks)
- P2P, sensor networks
- Grid computing (LHC, SETI, etc.)
- Social networks

# Service layers in a distributed system

- A distributed system organized as middleware
  - Middleware layer extends over multiple machines
  - Middleware masks heterogeneity and provides **transparency**






# Hiding complexity in a distributed system

- A distributed system will shield the programmer from some complex problems by “transparently” handling them
  - **Access** transparency: hide distribution (e.g., RMI)
  - **Location** transparency: hide resource location (e.g., URIs)
  - **Concurrency** transparency: hide concurrent operation on shared resources (e.g., transactions)
  - **Replication** transparency: hide service or resource redundancy (e.g., CDNs)
  - **Failure** transparency: hide failures of hardware or software components (e.g., transparent application failover)
  - And more...

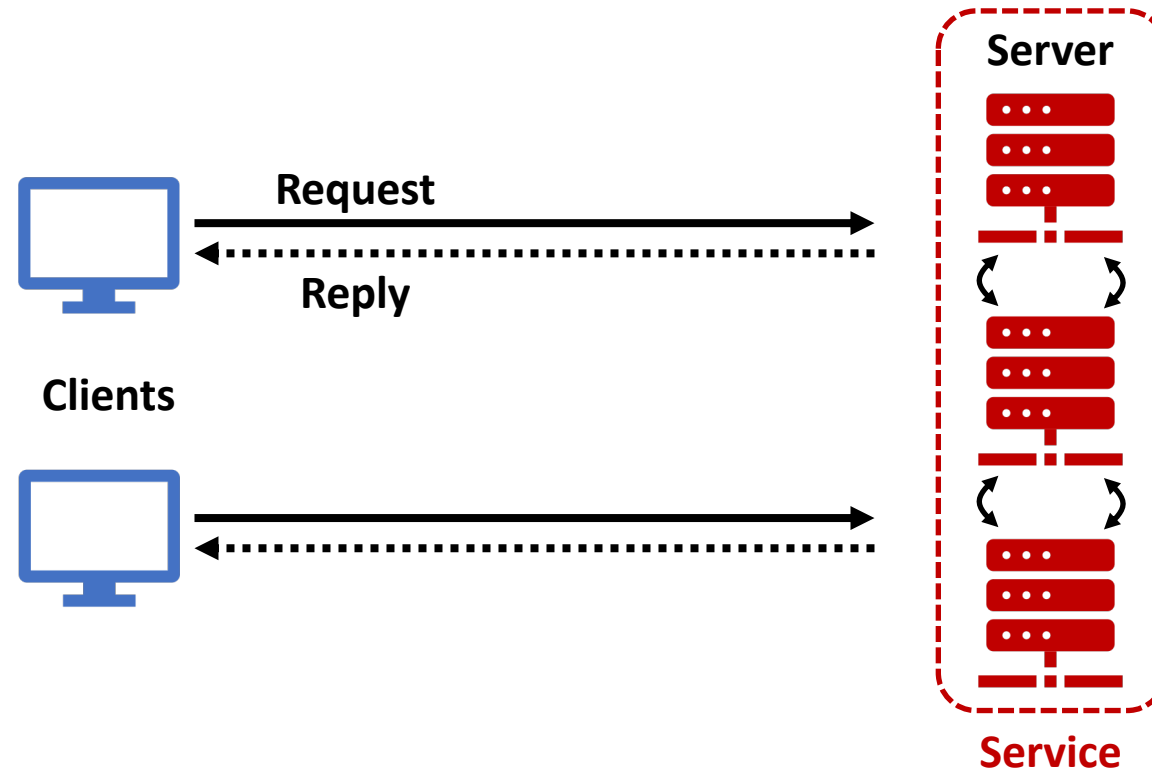
# Client/server architectures

- Clients invoke individual servers
- Servers may be clients of other servers
  - E.g., browser client of a Web server, in turn client of a file server, a database or a DNS




# Client/server architectures

- A service can be provided by multiple servers
  - E.g., Google, Amazon



# Peer architectures

- Multiple processes communicate with each other as peers
  - E.g., replicated service, file sharing networks, distributed whiteboard, internet routers



# Why distribute?

- Centralized architectures do not scale well
  - Centralized services (e.g., a single server for all users)
  - Centralized data (e.g., a single on-line telephone book)
  - Centralized algorithms (e.g., routing using complete information)
- Distribution enables **scaling**
  - Break components and spread parts across a distributed system
  - Use of asynchronous communication
  - Replicate critical components for high availability
  - Cache data (data replication)

# Why distribute?

- The **optimistic** view
  - Concurrency  $\Rightarrow$  speed (parallel processing, load balancing)
  - Partial failures  $\Rightarrow$  high availability
- The **pessimistic** view
  - Concurrency  $\Rightarrow$  incorrectness (interleavings)
  - Partial failures  $\Rightarrow$  incorrectness (inconsistent state)

# Some challenges of distributed computing

- Knowledge of a process
  - Local (identity, state, neighbours) vs. global (whole system)
- Network topology
  - Sparse topologies (e.g., ring, tree) vs. fully connected
- Degree of synchronization
  - Clock drift, message delays, processor speed
- Failure
  - Crash, omission, arbitrary
- Scalability
  - Performance not impaired by size, e.g.,  $O(\log N)$  vs.  $O(N)$

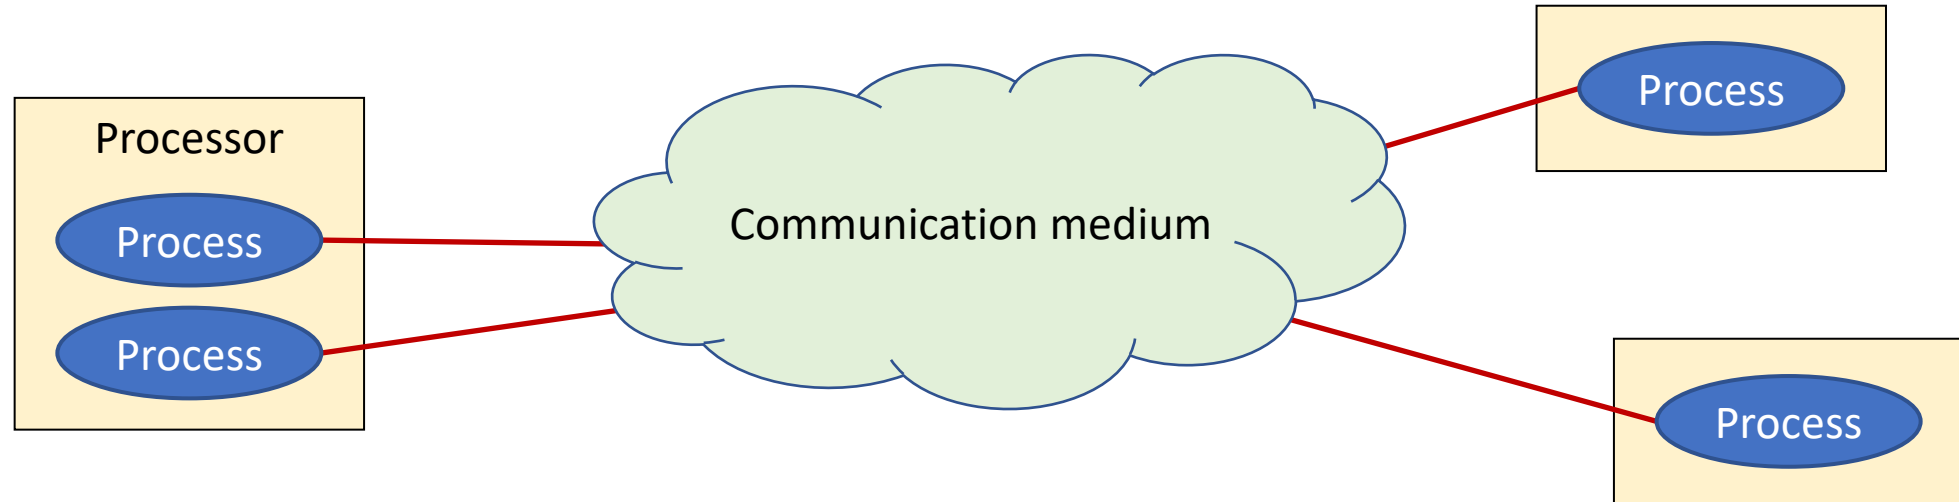
# Common subproblems

- Solving these challenges boils down to solving **subproblems**
  - Many applications revolve around a few common subproblems
  - Solving them helps having a good handle on system design
- Some classical examples
  - **Leader election**: elect one among a set of processes
  - **Mutual exclusion**: limit access to shared resources
  - **Time synchronization**: synchronize local clocks
  - **Global state**: collect all local states at a given time
  - **Multicast/broadcast**: send message to several processes
  - **Replica management**: keep state of replicas synchronized



# Distributed system models

- Processes communicating with one another



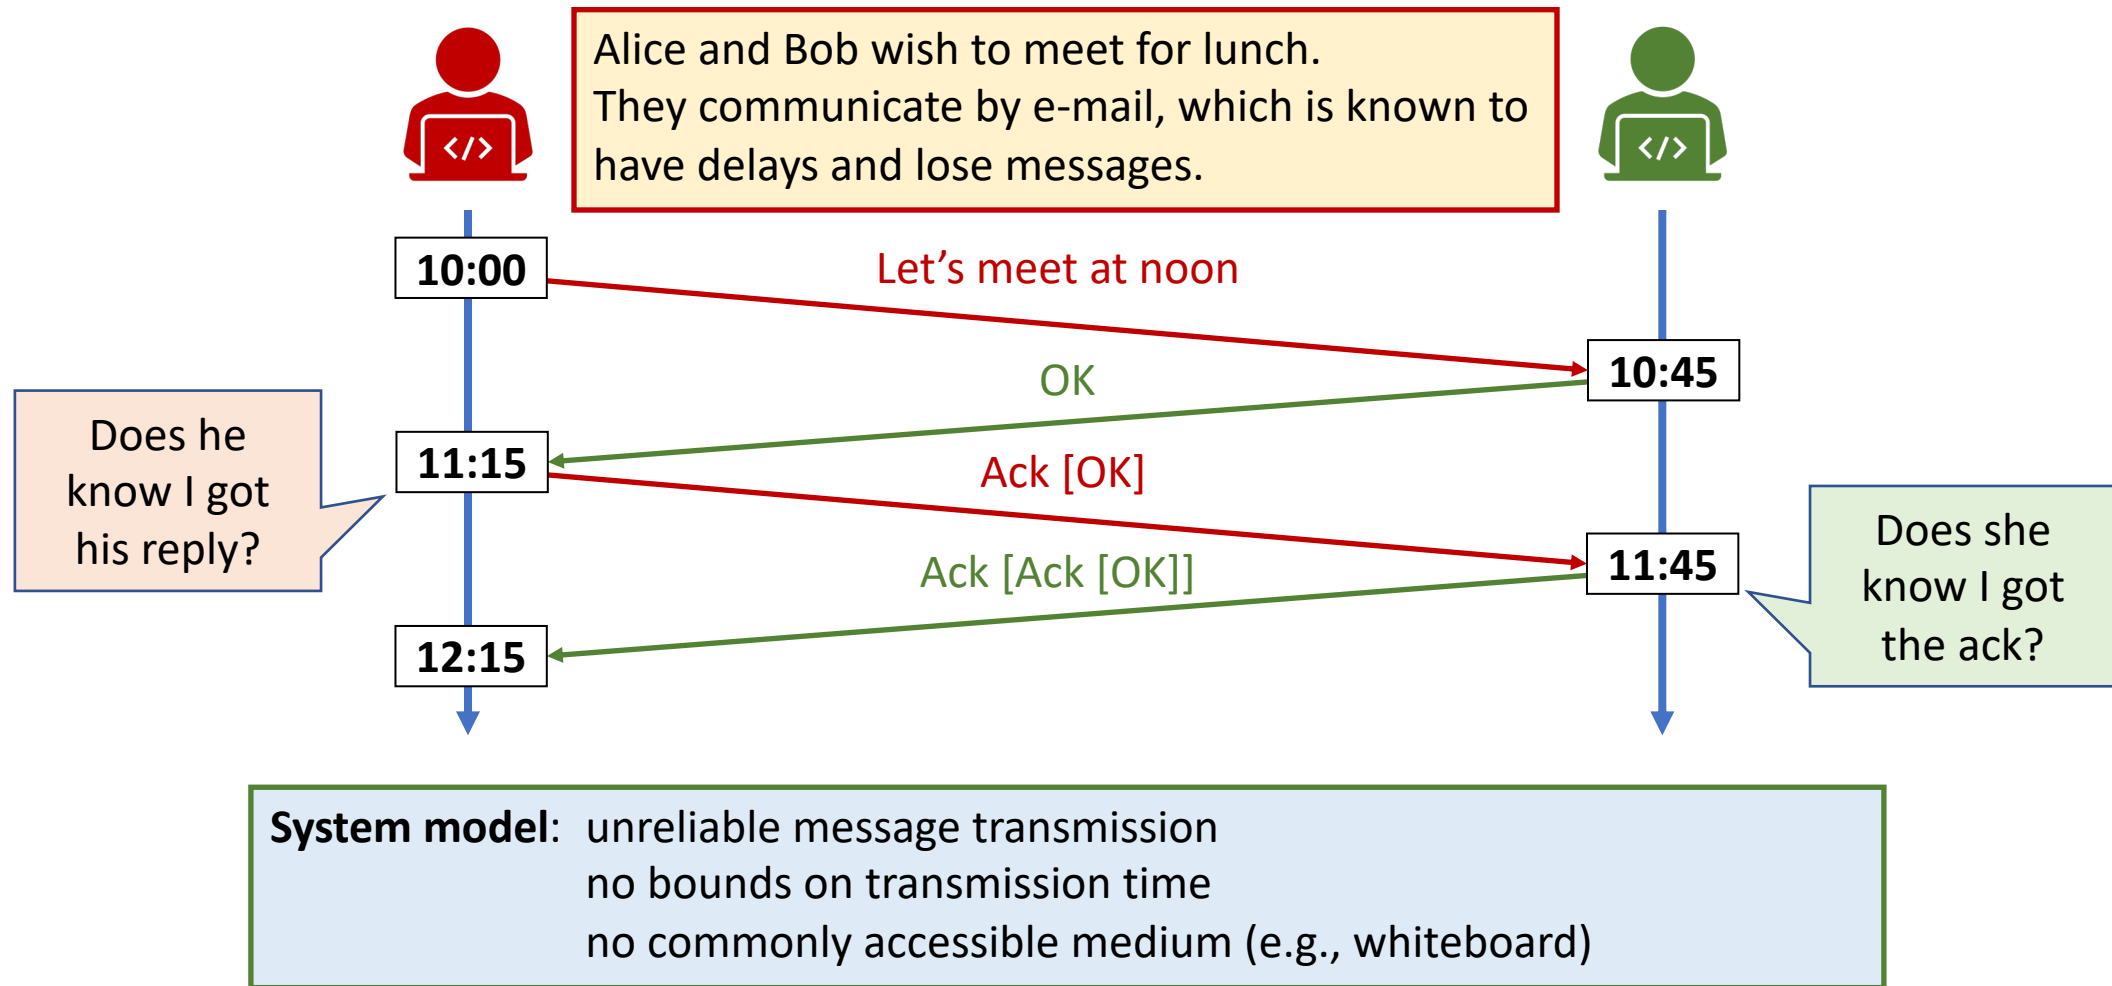
- What problems can one solve in a distributed system?
  - It depends on the **system model**: communication, timeliness, failure behaviour, etc.

# Solvability of distributed problems

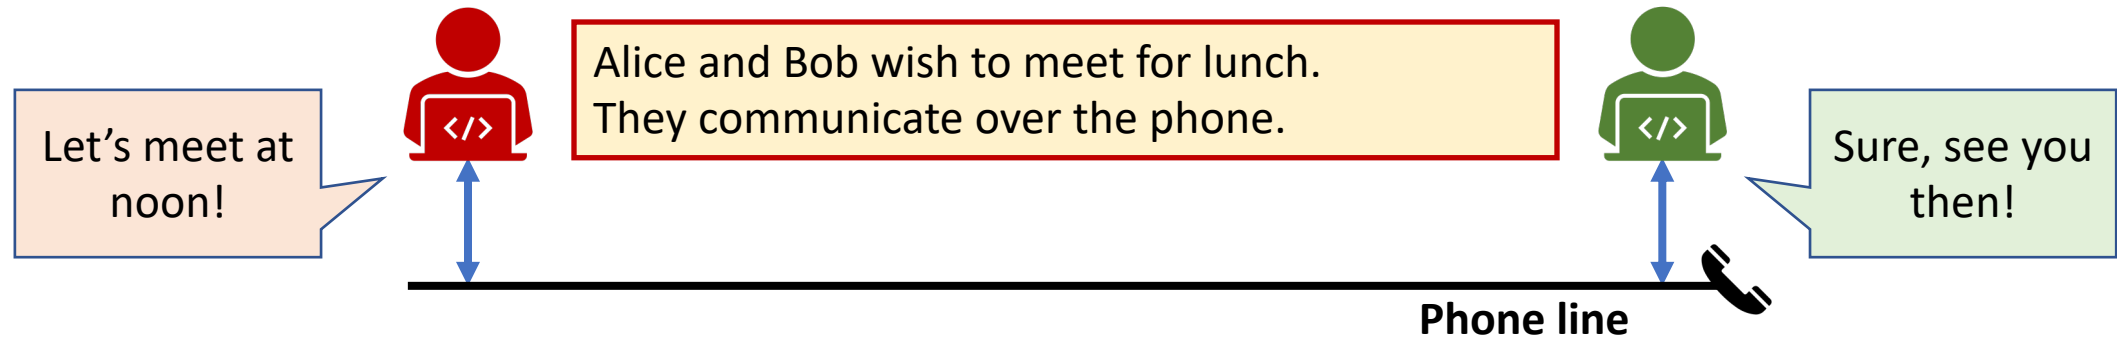
- Reliable broadcast
  - Ensure that a message sent to a group of processes is received by all or by none
- Atomic commit
  - Ensure that processes reach a common decision on whether to commit or abort a transaction
- Mutual exclusion
  - Ensure that only one process executes in a “critical section” at a time

“Solvability” and algorithms depends on the system model

# Example



# Example



**System model:** one party hears what the other says within a bounded delay  
*or*  
the existence of a problem is known within a bounded delay


# Some definitions

- A distributed system is made of a **finite set of processes** (generally denoted  $\Pi$  or  $P$ )
  - Each process models a **sequential** program
  - Processes are denoted by  $p_1, \dots, p_N$  or  $p, q, r$
  - Processes **communicate** with each other by exchanging messages

# Some definitions

- A process executes one step every tick of its local clock
  - A local computation (local event)
  - A message exchange (global event)
    - Sending a message to one or several processes
    - Receiving a message from one process
- A **history** of process  $p_i$  is an ordered series of events

$$h_i = \langle e_i^1, e_i^2, e_i^3, \dots \rangle$$



# Approach to distributed problems

- **Specifications**

*What is the problem?*

- **Assumptions**

*What is the system model?*

*What is the power of the adversary?*

- **Algorithms**

*How do we solve the problem?*

*How do we prove that the algorithm is correct?*

*At which cost?*

# Specifications

- A specification describes the problem in terms of safety and liveness properties
  - **Safety:** nothing bad ever happens
  - **Liveness:** something good eventually happens
- Any specification can be expressed in terms of safety and liveness properties



# Specifications

- Example of safety and liveness properties

*“Tell the truth!”*

**Safety:**      *“You shall not lie!”*

**Liveness:**    *“You have to say something”*

# Assumptions: system models

- But what is a model?
  - A collection of attributes and a set of rules that govern how these attributes interact


- Can a model be wrong?

*“A theory has only the alternative of being right or wrong.  
A model has a third possibility: it may be right, but irrelevant.”*

Manfred Eigen

# Good models

- Accurate models
  - Yield truth about the object of interest
- Tractable models
  - Analyzing them is actually possible



# Good models

- What to expect from a model?

- Feasibility

*What classes of problems can be solved?*  
(in a given model)

- Cost

*How expensive is the solution?*  
(for solvable problems)

# An example: coordination

- A coordination problem
  - Processes **p** and **q** communicate by sending and receiving messages on a bidirectional channel
  - Neither process can fail, but the channel may lose messages
  - Processes can execute one of two actions
  - Devise a protocol in which both processes take the same action, and neither takes both actions

# An example: coordination

- There is **no solution** to the problem!  
(in the given model)
- Proof (by contradiction)
  - Any protocol executes in rounds of message exchanges: first (say) **p** sends a message to **q**, then **q** sends a message to **p**, and so on
  - Let  $\Phi$  be the protocol that solves the problem using the fewest rounds
  - Assume w.l.o.g. that the last message is sent by **p**, and let it be **m**

# An example: coordination

- Proof (cont'd)
  - Observation #1: the action taken by **p** cannot depend on **m**, because its receipt could never be learned by **p** (it is the last message)
  - Observation #2: the action taken by **q** cannot depend on **m**, because **q** must make the same choice of action even if **m** is lost (due to a channel failure)

# An example: coordination

- Proof (cont'd)
  - Since the action chosen by **p** and **q** does not depend on **m**, it follows that **m** is not needed and so we can construct a new protocol in which one fewer message is sent...  
...a contradiction!




# Lessons learned


- What have we learned?
  - All protocols between two processes in this model are equivalent to a series of message exchanges
  - Actions taken by a process depend only on the sequence of messages it has received

# Assumptions: system models

- 3 dimensions to consider




# Communication medium



# Communication medium

- Point-to-point networks
  - Processes connected by a link communicate via **send**/**recv**
  - Links may be uni- or bi-directional
  - Non-blocking send primitive necessary for fault-tolerance
  - Often assume fully-connected network (links between each pair of processes)
  - Links are not necessarily direct physical connections



# Communication medium

- **Fair-lossy** (unreliable) point-to-point links (channels)

- **Fair-loss**

If a message  $m$  is sent **infinitely often** by  $p$  to  $q$ ,  
and **neither  $p$  nor  $q$  crashes**,  
then  $m$  is delivered **infinitely often** to  $q$

- **Finite duplication**

If a message  $m$  is sent a **finite number of times** by  $p$  to  $q$ ,  
then it is delivered a **finite number of times** to  $q$

- **No creation**

No message is delivered **unless** it was sent

# Communication medium

- **Reliable** point-to-point links (channels)

- **Validity**

If p sends m to q and q **does not fail**,  
then q **eventually receives** m from p

- **No duplication**

Process q receives messages m from p **at most once**

- **No creation**

Process q receives messages m from p  
only if p has **previously sent** m to q

# Failure types

- Link liveness failure: **message loss**

...message sent from **p** to correct process **q** never received by **q**...

- A link that **violates** (**satisfies**) specification is **faulty** (**correct**)

- Process liveness failure: **crash**

...process stops taking steps...

- A process that **violates** (**satisfies**) specification is **faulty** (**correct**)

**Crash-stop** model: a crashed process **never recovers** (it stops taking steps forever)

# Failure types

- Arbitrary failures (**Byzantine**)
  - Process/channel may send/transmit arbitrary messages at arbitrary times, or commit omissions
  - Process may take an incorrect step
- Performance failures
  - Process exceeds the bounds on the interval between two steps
  - Message transmission takes longer than the stated bound
- And more...




# Synchrony considerations

- **Synchronous** network model
  - Known upper bound on the time required for a process to execute a local step
  - Known upper bound on message transmission delay
  - Can assume that process have synchronized (within known bounds) physical clocks

# Synchrony considerations

- Consequences of synchronous model
  - Can use **timeouts** to detect process / link failures



- Can organize computation in **round**
  - Send messages to a set of processes  $\Pi$
  - Receive replies of that round from all processes in  $\Pi$  (failures are detected using timeouts)
  - Change state

# Synchrony considerations

- **Asynchronous** network model
  - **No bound** on the time required for a process to execute a local step (however, this time is finite)
  - **No bound** on the message transmission delay
  - **Cannot assume** the existence of perfectly or approximately **synchronized** physical clocks
- The most general model
  - An algorithm designed for asynchronous systems also works in synchronous systems




# Synchrony considerations

- Unfortunately, some very basic computational problems **cannot** be solved in *asynchronous systems*...  
in a fault-tolerant manner  
*and*  
with a deterministic algorithm
- Thus, for certain problems we have to resort to...  
synchronous (or partially synchronous) systems  
*or*  
randomized algorithm

# Common knowledge [Halpern et al.]

## The “muddy children” puzzle


- $N$  children play together,  
 $k$  get mud on their foreheads
- Each can see the mud on others  
but not on his forehead
- The father comes and says:  
*“At least one of you has mud  
on your forehead”*  
(a fact known by children)
- The father asks over and over:  
*“Does any of you know whether you  
have mud on your own forehead?”*
- All the children are perceptive,  
intelligent, truthful, and they  
answer simultaneously



# Common knowledge [Halpern et al.]


## The “muddy children” puzzle

- *N* children play together,  
*k* get mud on their foreheads
- Each can see the mud on others but not on his forehead
- The father comes and says:  
*“At least one of you has mud on your forehead”*  
(a fact known by children)
- The father asks over and over:  
*“Does any of you know whether you have mud on your own forehead?”*
- All the children are perceptive, intelligent, truthful, and they answer simultaneously



# Common knowledge [Halpern et al.]


- What's the use of the father's statement  
"At least one of you has a muddy forehead!"



- Alice, Carole and Daniel know that “at least one children has a muddy forehead”
- Bob does **not** know that “at least one children has a muddy forehead”

# Common knowledge [Halpern et al.]

- What's the use of the father's statement  
"At least one of you has a muddy forehead!"



- All know that "at least one children has a muddy forehead"
- Alice and Daniel know **that all know** that "at least one children has a muddy forehead..."  
...but Bob and Carole do **not** know **that all know** that "at least one children has a muddy forehead"




# Algorithms

- An algorithm is a solution to a problem
- Example: **reliable FIFO channels**
- Specification: **reliable channels + FIFO**


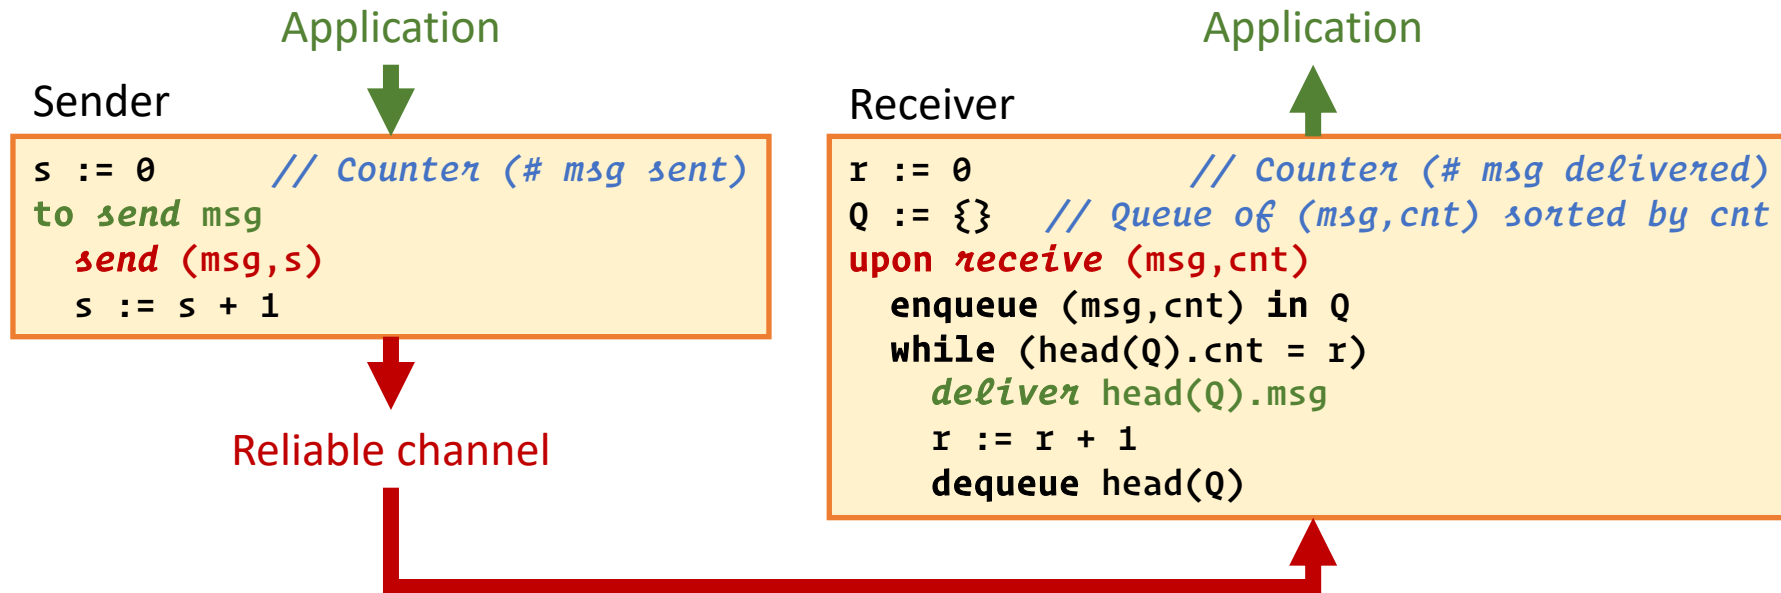
**FIFO:** if a process  $p$  sends  $m$  to  $q$  before sending  $m'$  to  $q$ ,  
then  $q$  does not deliver  $m'$  before  $m$

- Assumption: **reliable channels**



# Algorithms

**Proof:** follows from reliable channel assumption and code



# Summary

- To describe a distributed system, must specify **system model**
  - Communication, failures, synchrony
- Distributed computing problems can be specified by the means of **safety and liveness properties**...  
...and solved by **distributed algorithms**
- It is crucial to be clear and precise about these matters, as they affect whether
  - An algorithm works *in a given system*
  - A computational problem is solvable *in a given system*