

IN.5022 — Concurrent and Distributed Computing

Time and Order

Prof. P. Felber

pascal.felber@unine.ch

Agenda




- Time and order
- Clocks and physical time
- Logical clocks
- Vector clocks

Time: a major issue

- We casually use temporal concepts, mainly to measure time and order events
 - E.g., “upon timeout, rollback”, “once read lock is granted, acquire write lock” , “**p** suspect that **q** has failed”
- Used by many algorithms
 - E.g., distributed synchronization, maintain data consistency, authenticate requests, control concurrency
- In distributed systems, how can we relate local notion of time in a single process to a global notion of time?

Ordering of events?



Three notions of time

- Global clock
 - Time seen by external observer (wall clock time)
 - Hard to implement, limited temporal precision
- Local clocks of individual processes
 - Subject to skew and drift
 - Resynchronization is inaccurate
- Logical notion of time
 - Focus on relative ordering of events (occurred before)
 - No “real-time” clock

Physical time

Logical time

Time vs. ordering

- Time is often *wrongly* used to do ordering
 - E.g., make
- Time is useful for measuring intervals
 - E.g., performance analysis
- Ordering is useful for capturing temporal relationships
 - E.g., debugging (linearize observed set of events)
- How can we determine ordering in truly decentralized systems (many points of serialization)?

Computer clocks

- Each computer in a DS has its own internal clock
 - Used by local processes to obtain the current time value
 - Processes on different computers can timestamp events, but clocks on different computers may give different times
 - **Clock skew:** instantaneous differences between two clocks
 - Computer clocks “drift” from perfect time and their drift rates differ from one another
 - **Clock drift rate:** the relative amount that a computer clock differs from a perfect clock
- Even if clocks on all computers in a DS are set to the same time, their clocks will eventually vary quite significantly unless corrections are applied


Computer clocks

- To timestamp events, we can use the computer's clock
- At real time t , the OS reads the time on the computer's hardware clock $H_i(t)$
 - E.g., a 64-bit number giving nanoseconds since some base time
- It calculates the time on its software clock
$$C_i(t) = \alpha H_i(t) + \beta$$
 - In general, the clock is not completely accurate

How accurate should the clock resolution be?

Clock skew

- Computer clocks are generally *not* in perfect agreement (skew)
- Skew increases with drift
 - Ordinary quartz clocks drift by about 1 second every 11-12 days (10^{-6} s/s)
 - High precision quartz clocks drift rate is about 10^{-7} - 10^{-8} s/s



What happens to clocks when batteries become low?

- Computers must periodically synchronize their clocks!

Clock correctness

- A **hardware clock** H is said to be correct if its drift rate is within a bound $\rho > 0$ (e.g., 10^{-6} s/s)
- This means that the error in measuring the interval between real times t and t' ($t' > t$) is bounded
$$(1 - \rho)(t' - t) \leq H(t') - H(t) \leq (1 + \rho)(t' - t)$$
 - Bounded drift forbids jumps in time readings of hardware clocks

Clock correctness

- Weaker condition of monotonicity on **software clocks**
 $t' > t \Rightarrow C(t') > C(t)$ (clock value only ever increases)
 - E.g., required by Unix `make`
- We can achieve monotonicity with a hardware clock that runs fast by updating software clock at a slower rate
 - Adjust the values of α and β in $C_i(t) = \alpha H_i(t) + \beta$
- A **faulty clock** is one that does not obey its correctness condition
 - Crash failure:** a clock stops ticking
 - Arbitrary failure:** any other failure, e.g., jump back in time (Y2K)

Synchronizing physical clocks

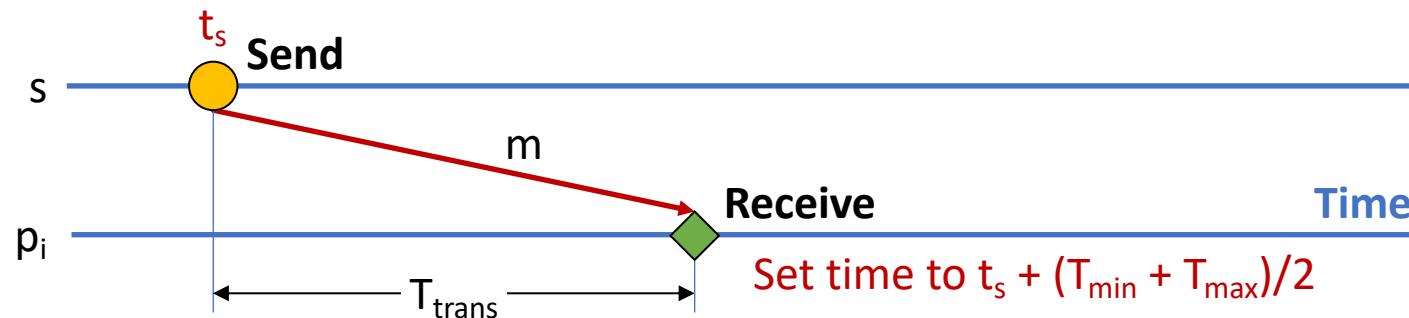
- External synchronization
 - A computer clock C_i is synchronized with an external authoritative time source S , so that
$$|S(t) - C_i(t)| < D \quad \text{for } i = 1, 2, \dots, N \text{ over an interval } I \text{ of real time}$$
 - The clock C_i is “*accurate to within the bound D* ”
- If two processes are synchronized externally within a bound D , then the reading of their clocks does not differ by more than twice D

Synchronizing physical clocks

- Internal synchronization
 - The clocks of each pair of computers are synchronized with one another so that
$$|C_i(t) - C_j(t)| < D \quad \text{for } i = 1, 2, \dots, N \text{ over an interval } I \text{ of real time}$$
 - The clocks C_i and C_j “*agree within the bound D* ”
- Internally synchronized clocks are not necessarily externally synchronized, as they may drift collectively
 - Often, this is not a problem...


Clocks in synchronous systems

- In a synchronous system
 - We know bounds on message transmission delay
$$T_{\min} \leq T_{\text{trans}} \leq T_{\max}$$
- External synchronization with time server s
 - Uncertainty $u = T_{\max} - T_{\min}$
 - Processes synchronized within $u/2$ with time server




Clocks in synchronous systems

- Internal synchronization
 - Process p_i sends its local time t to process p_j
 - Uncertainty $u = T_{\max} - T_{\min}$
 - Set clock to $t + (T_{\max} + T_{\min})/2 = t + T_{\text{med}} \Rightarrow \text{skew} \leq u/2$
 - With N processes, optimal precision of clocks C_i is $u(1 - 1/N)$




Clocks in asynchronous systems

- [Cristian 89] External synchronization with server s
 - Process p requests time in m_r and receives t in m_t from s
 - Let T_{RTT} be the RTT recorded by p and T_{min} the minimum transmission time
 - Process p sets its clock to $t + (T_{RTT}/2)$
 - Uncertainty is $T_{RTT} - 2T_{min}$ and accuracy is $\pm (T_{RTT}/2 - T_{min})$
 - Earliest time s puts t in message m_t is T_{min} after p sent m_r
 - Latest time was T_{min} before m_t arrived at p
 - Clock of s when m_t arrives is in $[t + T_{min}, t + T_{RTT} - T_{min}]$



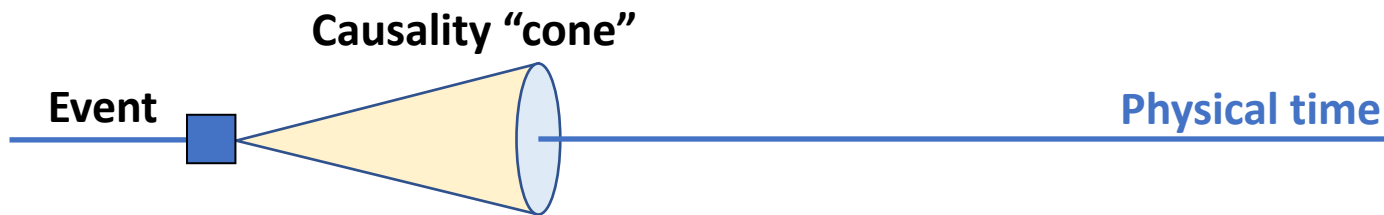
Clocks in asynchronous systems

- [Berkeley 89] Internal synchronization (group of computers)
 - A master **m** polls to collect clock values from others (slaves)
 - The master uses RTTs to estimate the slaves' clock values
 - It takes an average (eliminating dubious values)
 - It adjusts its clock and sends the required adjustment to the slaves
 - Better than sending the time, which depends on the RTT
 - If **m** fails, we can elect a new master (not in bounded time)




Logical time

- Alternative to synchronising physical clocks
- Event are uniquely ordered in any single process
 - _i: total order defined by the order in which p_i observes events
- Distributed events are ordered according to causality
 - When a message m is sent, $\text{send}(m)$ occurs before $\text{receive}(m)$
- Note: events propagate at a finite speed



“Potential” causality

- “Happened-before” relation (\rightarrow)
 - [HB1]: if \exists process $p_i : e \rightarrow_i e'$, then $e \rightarrow e'$
 - [HB2]: \forall message m , $\text{send}(m) \rightarrow \text{receive}(m)$
 - [HB3]: if $e \rightarrow e'$ and $e' \rightarrow e''$, then $e \rightarrow e''$ (transitive)
- Partial order
 - For some events, we do not know which one happened first
- Concurrent events
 - If $e \nrightarrow f$ and $f \nrightarrow e$, then $e \parallel f$



Logical clocks definition

- A logical clock **C** is a mapping from the set of states **S** to **N** (natural numbers) with the following constraint

$$\forall s, t \in S : s < t \vee s \rightsquigarrow t \Rightarrow C(s) < C(t)$$

$<$: locally precedes

\rightsquigarrow : remotely precedes

Logical clocks algorithm

- Introduced by Leslie Lamport in 1978
 - Orders events globally according to the \rightarrow relation
- L_i : logical clock (counter) used by process p_i to apply logical (“Lamport”) timestamp $L(e)$ to event e

[LC1]


L_i incremented before each event at p_i : $L_i := L_i + 1$

[LC2]

- a) When process p_i sends message m , it piggybacks on m the value $t = L_i$
- b) On receiving (m, t) , p_j computes $L_j := \max(L_j, t)$ and then applies LC1 before timestamping $\text{receive}(m)$

$$e \rightarrow e' \Rightarrow L(e) < L(e')$$

Logical clocks example




Logical clocks and total order

- Some events have the same Lamport timestamp
 - E.g., $L(a) = L(e)$
- Break ties by using processes ranks
 - Local timestamp T_i at process p_i becomes global timestamp (T_i, i)
 $(T_i, i) < (T_j, j) \Leftrightarrow T_i < T_j \vee (T_i = T_j \wedge i < j)$
 - Global timestamps form a total order

Is that good enough?

Logical clocks limitations

- Are events **a** and **c** ordered?
 - Yes, and $L(a) < L(c)$
- Are events **e** and **c** ordered?
 - No, **but** $L(e) < L(c)$



Logical clocks limitations

- For some pair of events, we do not know which happened first (partial ordering)
 - When ordering is unknown, an arbitrary order is chosen
 - We cannot find true dependencies by looking at ordering
 - Timestamps do not distinguish between causally and arbitrarily ordered events

Vector clocks

- Introduced (independently) by Colin Fidge and Friedemann Mattern in 1988
- Overcome main shortcoming of Lamport's clocks
$$L(e) < L(e') \not\Rightarrow e \rightarrow e'$$
- Preserve partial ordering information
- If ordering of events is unknown, leave them unordered (incomparable events)
 - Easier to detect race conditions

Vector clocks concepts

- Each process has an array of logical clocks
 - One clock per process in the system
 - Vector of last known timestamps

$$V_i = (t_{p_0}, t_{p_1}, t_{p_2}, \dots, t_{p_N})$$

- Every event is given a timestamp vector by the process to which it belongs
- The ordering, or lack thereof, of two events can be determined by comparing their timestamps

Vector clocks definition

- A vector clock V is a mapping from the set of states S to \mathbb{N}^k (vector of natural numbers) with the following constraint

$$\forall s, t \in S : s \rightarrow t \Rightarrow V(s) < V(t)$$

\rightarrow is a partial order, thus $<$ must also be a partial order

Vector clocks algorithm

- V_i : vector clock used by process p_i to timestamp events

[VC1]

Initially, $V_i[j] = 0$ for $i, j = 1, 2, \dots, N$

[VC2]

Before timestamping e , p_i sets $V_i[i] := V_i[i] + 1$


[VC3]

When process p_i sends message m , it piggybacks on m the value $t = V_i$

[VC4]

On receiving (m, t) , process p_i computes $V_i[j] := \max(V_i[j], t[j])$ for $j = 1, 2, \dots, N$ and then applies VC2 before timestamping the event $\text{receive}(m)$

Vector clocks example




Vector clocks properties

- Interpretation
 - $V_i[i]$: number of events that p_i has timestamped
 - $V_i[j]$ ($j \neq i$): number of events occurred at p_j that p_i has potentially been affected by
- Comparing vector clocks
 - $V = V' \Leftrightarrow V[j] = V'[j]$ for $j = 1, 2, \dots, N$
 - $V \leq V' \Leftrightarrow V[j] \leq V'[j]$ for $j = 1, 2, \dots, N$
 - $V < V' \Leftrightarrow V \leq V' \wedge V \neq V'$
 - E.g., $(2,1,0,4) < (2,3,0,4)$

Vector clocks benefits

- Are events **a** and **c** ordered?
 - Yes, because $V(a) < V(c)$
- Are events **e** and **c** ordered?
 - No, because $V(e) \not\leq V(c) \wedge V(c) \not\leq V(e)$



Vector clocks pros and cons

- 😊 We have $e \rightarrow e' \Leftrightarrow V(e) < V(e')$
 - Can tell whether e “happened before” e' from vector clocks
- 😊 If $V(e) \not\leq V(e')$ and $V(e') \not\leq V(e)$, then $e \parallel e'$
 - Events are concurrent when vector clocks are not comparable
 - E.g., $(2,1,0,4) \parallel (2,3,0,2)$
- 😞 Requires static notion of system membership
 - Processes must agree on the number of entries in vectors
 - Vector clocks are useful in systems that deal with membership, e.g., group communication
 - There are techniques to deal with dynamic group membership

Example: atomic broadcast


- A process sends a message atomically to all other processes
 - [Agreement]: if some correct process delivers **m**, then all correct processes deliver **m**
 - [Ordering]: no two correct processes deliver any two messages in different orders
 - [Termination]: if a correct process broadcasts **m**, then all correct processes eventually deliver **m**

How can we implement atomic broadcast using logical time?

System model: reliable channels, asynchronous system, no failures

Atomic broadcast protocol

- Sender adds timestamp in the broadcast
- Receiver waits for full set of messages
 - Orders messages by logical timestamp
 - Breaks ties using sender identifiers
 - Delivers messages in this order
 - Picks new timestamp greater than all seen
- How do we know if we have a full set?
 - Rely upon “membership” to wait for (sets of) messages from *all members*
 - System runs in rounds
 - Send “null” message if nothing to send



Deliver m (timestamp 2)
before
 m' (timestamp 3)

Interpretation of logical time

- The relation “**a** happened before **b**” means that information can flow from **a** to **b**
- The relation “**a** is concurrent with **b**” means that no information can flow between **a** to **b**
 - Many events can be concurrent with a given event
 - Logical time cannot help detect “simultaneous” events
 - “Real-time” clocks cannot help either, because of limited precision and communication latencies
 - Useful only for “coarse-grain” applications

Things to remember

- Accurate timekeeping is important in a distributed system
 - Algorithms synchronize clocks despite their drift and the variability of message delays
- Time is a tool, typically used to put events in some sort of order
 - E.g., order updates on replicated data
- Often physical time is not necessary and logical time can be used instead
- Logical and vector clocks provide a *partial* order
 - Can be extended to a total order, e.g., by adding clock time or process identifiers to break ties